STEP 1 — Install & Import Libraries

```
# Install required libraries (run once in notebook)
!pip install torch transformers datasets peft accelerate

# Import libraries
import torch
from transformers import AutoModelForSequenceClassification, AutoTokenizer
from datasets import load_dataset
from peft import LoraConfig, get_peft_model
```

```
Requirement already satisfied: torch in /usr/local/lib/python3.12/dist-packages (2.10.0+cpu)
Requirement already satisfied: transformers in /usr/local/lib/python3.12/dist-packages (5.0.0)
Requirement already satisfied: datasets in /usr/local/lib/python3.12/dist-packages (4.0.0)
Requirement already satisfied: peft in /usr/local/lib/python3.12/dist-packages (0.18.1)
Requirement already satisfied: accelerate in /usr/local/lib/python3.12/dist-packages (1.13.0)
Requirement already satisfied: filelock in /usr/local/lib/python3.12/dist-packages (from torch) (3.25.2)
Requirement already satisfied: typing-extensions>=4.10.0 in /usr/local/lib/python3.12/dist-packages (from torch) (4.15.0)
Requirement already satisfied: setuptools in /usr/local/lib/python3.12/dist-packages (from torch) (75.2.0)
Requirement already satisfied: sympy>=1.13.3 in /usr/local/lib/python3.12/dist-packages (from torch) (1.14.0)
Requirement already satisfied: networkx>=2.5.1 in /usr/local/lib/python3.12/dist-packages (from torch) (3.6.1)
Requirement already satisfied: jinja2 in /usr/local/lib/python3.12/dist-packages (from torch) (3.1.6)
Requirement already satisfied: fsspec>=0.8.5 in /usr/local/lib/python3.12/dist-packages (from torch) (2025.3.0)
Requirement already satisfied: huggingface-hub<2.0,>=1.3.0 in /usr/local/lib/python3.12/dist-packages (from transformers)
Requirement already satisfied: numpy>=1.17 in /usr/local/lib/python3.12/dist-packages (from transformers) (2.0.2)
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.12/dist-packages (from transformers) (26.0)
Requirement already satisfied: pyyaml>=5.1 in /usr/local/lib/python3.12/dist-packages (from transformers) (6.0.3)
Requirement already satisfied: regex!=2019.12.17 in /usr/local/lib/python3.12/dist-packages (from transformers) (2025.11.3
Requirement already satisfied: tokenizers<=0.23.0,>=0.22.0 in /usr/local/lib/python3.12/dist-packages (from transformers)
Requirement already satisfied: typer-slim in /usr/local/lib/python3.12/dist-packages (from transformers) (0.24.0)
Requirement already satisfied: safetensors>=0.4.3 in /usr/local/lib/python3.12/dist-packages (from transformers) (0.7.0)
Requirement already satisfied: tqdm>=4.27 in /usr/local/lib/python3.12/dist-packages (from transformers) (4.67.3)
Requirement already satisfied: pyarrow>=15.0.0 in /usr/local/lib/python3.12/dist-packages (from datasets) (18.1.0)
Requirement already satisfied: dill<0.3.9,>=0.3.0 in /usr/local/lib/python3.12/dist-packages (from datasets) (0.3.8)
Requirement already satisfied: pandas in /usr/local/lib/python3.12/dist-packages (from datasets) (2.2.2)
Requirement already satisfied: requests>=2.32.2 in /usr/local/lib/python3.12/dist-packages (from datasets) (2.32.4)
Requirement already satisfied: xxhash in /usr/local/lib/python3.12/dist-packages (from datasets) (3.6.0)
Requirement already satisfied: multiprocess<0.70.17 in /usr/local/lib/python3.12/dist-packages (from datasets) (0.70.16)
Requirement already satisfied: psutil in /usr/local/lib/python3.12/dist-packages (from peft) (5.9.5)
Requirement already satisfied: aiohttp!=4.0.0a0,!=4.0.0a1 in /usr/local/lib/python3.12/dist-packages (from fsspec[http]<=2
Requirement already satisfied: hf-xet<2.0.0,>=1.4.2 in /usr/local/lib/python3.12/dist-packages (from huggingface-hub<2.0,>
Requirement already satisfied: httpx<1,>=0.23.0 in /usr/local/lib/python3.12/dist-packages (from huggingface-hub<2.0,>=1.3
Requirement already satisfied: typer in /usr/local/lib/python3.12/dist-packages (from huggingface-hub<2.0,>=1.3.0->transfo
Requirement already satisfied: charset_normalizer<4,>=2 in /usr/local/lib/python3.12/dist-packages (from requests>=2.32.2-
Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.12/dist-packages (from requests>=2.32.2->datasets) (
Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/local/lib/python3.12/dist-packages (from requests>=2.32.2->datas
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.12/dist-packages (from requests>=2.32.2->datas
Requirement already satisfied: mpmath<1.4,>=1.1.0 in /usr/local/lib/python3.12/dist-packages (from sympy>=1.13.3->torch) (
Requirement already satisfied: MarkupSafe>=2.0 in /usr/local/lib/python3.12/dist-packages (from jinja2->torch) (3.0.3)
Requirement already satisfied: python-dateutil>=2.8.2 in /usr/local/lib/python3.12/dist-packages (from pandas->datasets) (
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.12/dist-packages (from pandas->datasets) (2025.2)
Requirement already satisfied: tzdata>=2022.7 in /usr/local/lib/python3.12/dist-packages (from pandas->datasets) (2025.3)
Requirement already satisfied: aiohappyeyeballs>=2.5.0 in /usr/local/lib/python3.12/dist-packages (from aiohttp!=4.0.0a0,!
Requirement already satisfied: aiosignal>=1.4.0 in /usr/local/lib/python3.12/dist-packages (from aiohttp!=4.0.0a0,!=4.0.0a
Requirement already satisfied: attrs>=17.3.0 in /usr/local/lib/python3.12/dist-packages (from aiohttp!=4.0.0a0,!=4.0.0a1->
Requirement already satisfied: frozenlist>=1.1.1 in /usr/local/lib/python3.12/dist-packages (from aiohttp!=4.0.0a0,!=4.0.0
Requirement already satisfied: multidict<7.0,>=4.5 in /usr/local/lib/python3.12/dist-packages (from aiohttp!=4.0.0a0,!=4.0
Requirement already satisfied: propcache>=0.2.0 in /usr/local/lib/python3.12/dist-packages (from aiohttp!=4.0.0a0,!=4.0.0a
Requirement already satisfied: yarl<2.0,>=1.17.0 in /usr/local/lib/python3.12/dist-packages (from aiohttp!=4.0.0a0,!=4.0.0
Requirement already satisfied: anyio in /usr/local/lib/python3.12/dist-packages (from httpx<1,>=0.23.0->huggingface-hub<2.
Requirement already satisfied: httpcore==1.* in /usr/local/lib/python3.12/dist-packages (from httpx<1,>=0.23.0->huggingfac
Requirement already satisfied: h11>=0.16 in /usr/local/lib/python3.12/dist-packages (from httpcore==1.*->httpx<1,>=0.23.0-
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.12/dist-packages (from python-dateutil>=2.8.2->pandas->d
Requirement already satisfied: click>=8.2.1 in /usr/local/lib/python3.12/dist-packages (from typer->huggingface-hub<2.0,>=
Requirement already satisfied: shellingham>=1.3.0 in /usr/local/lib/python3.12/dist-packages (from typer->huggingface-hub<
Requirement already satisfied: rich>=12.3.0 in /usr/local/lib/python3.12/dist-packages (from typer->huggingface-hub<2.0,>=
Requirement already satisfied: annotated-doc>=0.0.2 in /usr/local/lib/python3.12/dist-packages (from typer->huggingface-hu
Requirement already satisfied: markdown-it-py>=2.2.0 in /usr/local/lib/python3.12/dist-packages (from rich>=12.3.0->typer-
```

STEP 2 — Load Pre-trained Model & Tokenizer

```
model_name = "distilbert-base-uncased"

# Load tokenizer
tokenizer = AutoTokenizer.from_pretrained(model_name)

# Load model
model = AutoModelForSequenceClassification.from_pretrained(
    model_name,
    num_labels=2
```

```
)

print("Model and tokenizer loaded successfully!")
```

```
/usr/local/lib/python3.12/dist-packages/huggingface_hub/utils/_auth.py:94: UserWarning:
The secret `HF_TOKEN` does not exist in your Colab secrets.
To authenticate with the Hugging Face Hub, create a token in your settings tab (https://huggingface.co/settings/tokens), set
You will be able to reuse this secret in all of your notebooks.
Please note that authentication is recommended but still optional to access public models or datasets.
  warnings.warn(
config.json: 100% 483/483 [00:00<00:00, 25.4kB/s]
Warning: You are sending unauthenticated requests to the HF Hub. Please set a HF_TOKEN to enable higher rate limits and fast
WARNING:huggingface_hub.utils._http:Warning: You are sending unauthenticated requests to the HF Hub. Please set a HF_TOKEN t
tokenizer_config.json: 100% 48.0/48.0 [00:00<00:00, 1.22kB/s]
vocab.txt: 232k/? [00:00<00:00, 4.71MB/s]
tokenizer.json: 466k/? [00:00<00:00, 11.2MB/s]
model.safetensors: 100% 268M/268M [00:03<00:00, 460MB/s]
Loading weights: 100% 100/100 [00:00<00:00, 211.07it/s, Materializing param=distilbert.transformer.layer.5.sa_layer_norm.weight]
DistilBertForSequenceClassification LOAD REPORT from: distilbert-base-uncased
Key                     | Status     |
------------------------+------------+-
vocab_layer_norm.bias   | UNEXPECTED |
vocab_transform.weight  | UNEXPECTED |
vocab_projector.bias    | UNEXPECTED |
vocab_transform.bias    | UNEXPECTED |
vocab_layer_norm.weight | UNEXPECTED |
classifier.weight       | MISSING    |
classifier.bias         | MISSING    |
pre_classifier.bias     | MISSING    |
pre_classifier.weight   | MISSING    |

Notes:
- UNEXPECTED    :can be ignored when loading from different task/architecture; not ok if you expect identical arch.
- MISSING   :those params were newly initialized because missing from the checkpoint. Consider training on your downstream t
Model and tokenizer loaded successfully!
```

STEP 3 — Load & Preprocess Dataset

```
# Load dataset (IMDB sentiment dataset)
dataset = load_dataset("imdb")

# Tokenization function
def tokenize_function(example):
    return tokenizer(example["text"], truncation=True, padding="max_length")

# Apply tokenization
tokenized_dataset = dataset.map(tokenize_function, batched=True)

# Select smaller subset (as per lab requirement)
train_dataset = tokenized_dataset["train"].shuffle(seed=42).select(range(500))
test_dataset = tokenized_dataset["test"].shuffle(seed=42).select(range(200))

print("Dataset loaded and tokenized!")
```

```
README.md: 7.81k/? [00:00<00:00, 151kB/s]
plain_text/train-00000-of-00001.parquet: 100% 21.0M/21.0M [00:00<00:00, 84.7MB/s]
plain_text/test-00000-of-00001.parquet: 100% 20.5M/20.5M [00:00<00:00, 102MB/s]
plain_text/unsupervised-00000-of-00001.p(…): 100% 42.0M/42.0M [00:01<00:00, 68.6MB/s]
Generating train split: 100% 25000/25000 [00:01<00:00, 24302.84 examples/s]
Generating test split: 100% 25000/25000 [00:01<00:00, 21378.24 examples/s]
Generating unsupervised split: 100% 50000/50000 [00:02<00:00, 29306.16 examples/s]
Map: 100% 25000/25000 [01:02<00:00, 352.46 examples/s]
Map: 100% 25000/25000 [00:42<00:00, 842.50 examples/s]
Map: 100% 50000/50000 [01:08<00:00, 657.56 examples/s]
Dataset loaded and tokenized!
```

STEP 4 — Configure LoRA Adapter

```
lora_config = LoraConfig(
    r=8,                        # rank
    lora_alpha=16,              # scaling factor
    target_modules=["q_lin", "v_lin"],  # attention layers
    lora_dropout=0.1,
    bias="none",
    task_type="SEQ_CLS"
)

print("LoRA configuration created!")
```

```
LoRA configuration created!
```

STEP 5 — Apply PEFT Model

```
# Attach LoRA to model
peft_model = get_peft_model(model, lora_config)

# Print trainable parameters
peft_model.print_trainable_parameters()
```

```
WARNING:torchao.kernel.intmm:Warning: Detected no triton, on systems without Triton certain kernels will not work
trainable params: 739,586 || all params: 67,694,596 || trainable%: 1.0925
```

STEP 6 — Train the Model

```
from transformers import TrainingArguments, Trainer

training_args = TrainingArguments(
    output_dir="./results",
    learning_rate=2e-4,
    per_device_train_batch_size=8,
    num_train_epochs=1,
    logging_steps=10,
    save_strategy="no"
)

trainer = Trainer(
    model=peft_model,
    args=training_args,
    train_dataset=train_dataset
)

# Train
trainer.train()
```

```
/usr/local/lib/python3.12/dist-packages/torch/utils/data/dataloader.py:775: UserWarning: 'pin_memory' argument is set as tru
  super().__init__(loader)
```

[57/63 18:57 < 02:04, 0.05 it/s, Epoch 0.89/1]

Step	Training Loss
10	0.705426
20	0.664024
30	0.673471
40	0.637650
50	0.635734

STEP 7 — Evaluate & Predict

```
# Evaluate model
results = trainer.evaluate()
print("Evaluation Results:", results)

# Test prediction
text = "The movie was fantastic and inspiring!"
inputs = tokenizer(text, return_tensors="pt")

outputs = peft_model(**inputs)

prediction = torch.argmax(outputs.logits, dim=1)
```

```
print("Predicted label:", prediction.item())
```

STEP 8 — Comparison: PEFT vs Full Fine-Tuning

Parameter-Efficient Fine-Tuning (PEFT) and full fine-tuning differ significantly in terms of computational cost, efficiency, and scalability. In full fine-tuning, all the parameters of the pre-trained model are updated during training. This requires a large amount of GPU memory and computational power, especially for large transformer models with millions or billions of parameters. In contrast, PEFT updates only a small subset of parameters, such as adapter layers introduced through methods like LoRA.

One of the main advantages of PEFT is reduced memory usage. Since most of the model parameters are frozen, only a few additional parameters are trained, which drastically lowers memory requirements. This makes it feasible to fine-tune large models even on limited hardware like standard GPUs or cloud notebooks.

Training time is also significantly reduced in PEFT. Because fewer parameters are updated, the optimization process becomes faster compared to full fine-tuning. This allows quicker experimentation and iteration during development.

In terms of accuracy, PEFT often achieves performance comparable to full fine-tuning, especially for many NLP tasks such as sentiment analysis and classification. While there might be a slight drop in performance in some cases, the trade-off is usually acceptable given the efficiency gains.

Another important difference is parameter storage. Full fine-tuning requires storing an entirely new version of the model, which consumes large disk space. PEFT, however, only stores the small adapter weights, making it more storage-efficient and easier to deploy.

Overall, PEFT provides a practical and scalable solution for adapting large language models, while full fine-tuning is more resource-intensive and less efficient for most real-world applications.

STEP 9 — LAB REPORT

🔷 1. Objective

The objective of this lab is to understand and implement Parameter-Efficient Fine-Tuning (PEFT) using the LoRA technique. The goal is to efficiently fine-tune a pre-trained transformer model while minimizing computational cost and memory usage.

🔷 2. Dataset Description

The dataset used in this lab is the IMDB movie review dataset for sentiment analysis. It consists of text reviews labeled as positive or negative. The dataset contains natural language sentences with varying lengths and informal expressions. A subset of 500 training samples and 200 testing samples was used. The dataset was preprocessed by tokenizing the text using a transformer tokenizer, converting it into numerical format suitable for model training.

🔷 3. Installation and Setup

The following libraries were installed and used:

PyTorch Hugging Face Transformers Datasets library PEFT library Accelerate

Installation command:

pip install torch transformers datasets peft accelerate

The environment was set up using Python and executed in Jupyter Notebook/Google Colab.

🔷 4. LoRA Configuration Details

LoRA (Low-Rank Adaptation) was applied to the transformer model using the following configuration:

Rank (r): 8 Alpha: 16 Dropout: 0.1 Target modules: Query and Value projection layers

LoRA works by introducing low-rank matrices into attention layers, reducing the number of trainable parameters while maintaining performance.

🔷 5. Training Results

The model was trained for 1 epoch using a small dataset. The training process was efficient due to reduced parameter updates. Logging was performed at regular intervals. The model showed good learning behavior even with limited training data.

🔷 6. Performance Analysis

The trained model was evaluated on a test dataset. It was able to correctly classify sentiment in most cases. The predictions were generated by passing input text through the tokenizer and model. The performance was satisfactory considering the small dataset and limited training time.

🔷 7. Comparison with Traditional Fine-Tuning

Compared to full fine-tuning, PEFT required significantly less memory and training time. Only a small fraction of parameters were updated, making it computationally efficient. While full fine-tuning may achieve slightly higher accuracy, PEFT provides a better balance between performance and efficiency. It is especially useful when working with large models or limited hardware.

🔷 8. Conclusion

In this lab, we successfully implemented Parameter-Efficient Fine-Tuning using LoRA. The approach demonstrated that it is possible to adapt large transformer models efficiently without updating all parameters. PEFT is a powerful technique for modern NLP tasks, enabling faster training, lower memory usage, and easier deployment.